

COMP3308/COMP3608, Lecture 3a

ARTIFICIAL INTELLIGENCE

A* Algorithm


Reference: Russell and Norvig, ch. 3

Outline

- **A* search algorithm**
- **How to invent admissible heuristics**

A* Search


- UCS minimizes the cost so far $g(n)$
- GS minimizes the estimated cost to the goal $h(n)$
- A* combines UCS and GS
- Evaluation function: $f(n)=g(n)+h(n)$
 - $g(n)$ = cost so far to reach n
 - $h(n)$ = estimated cost from n to the goal
 - $f(n)$ = estimated total cost of path through n to the goal




A* Search for Romania Example



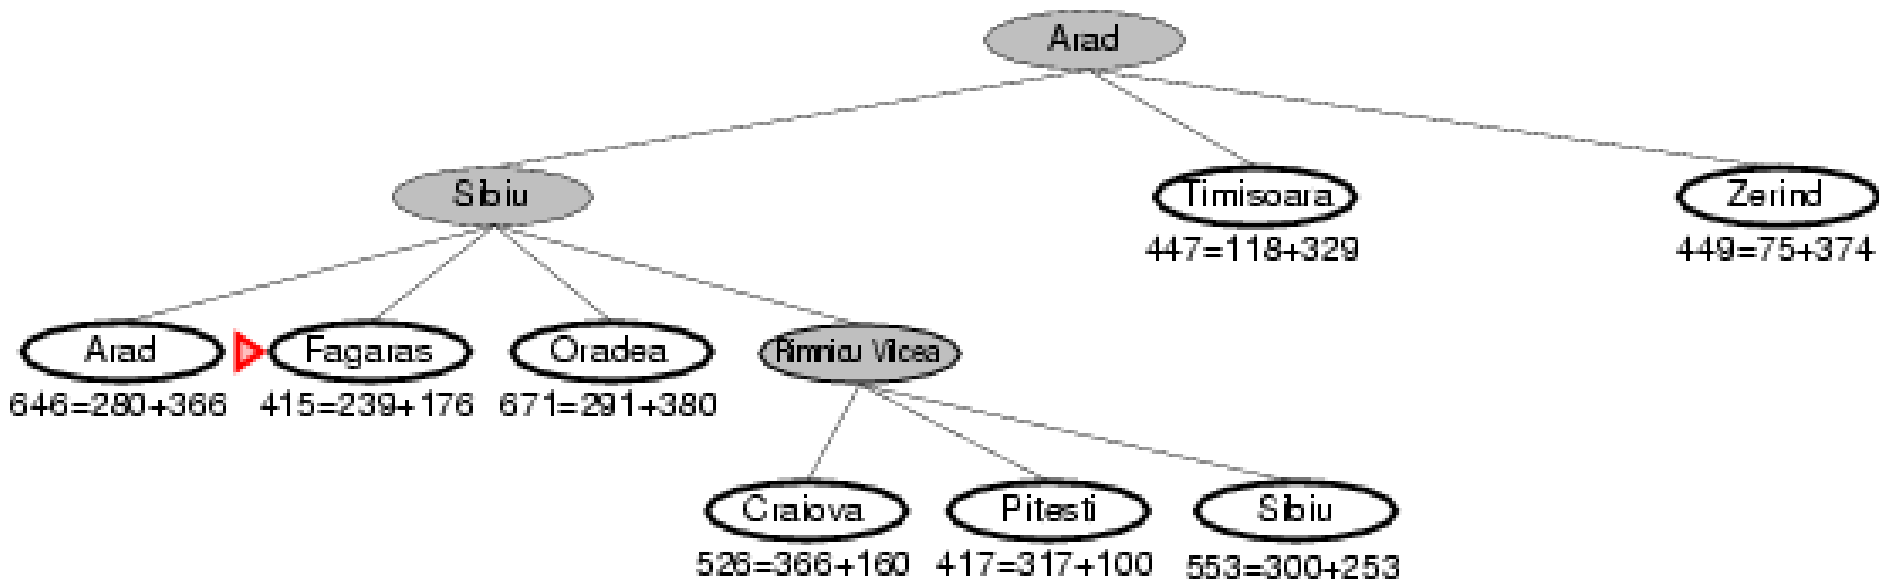
A* Search for Romania Example



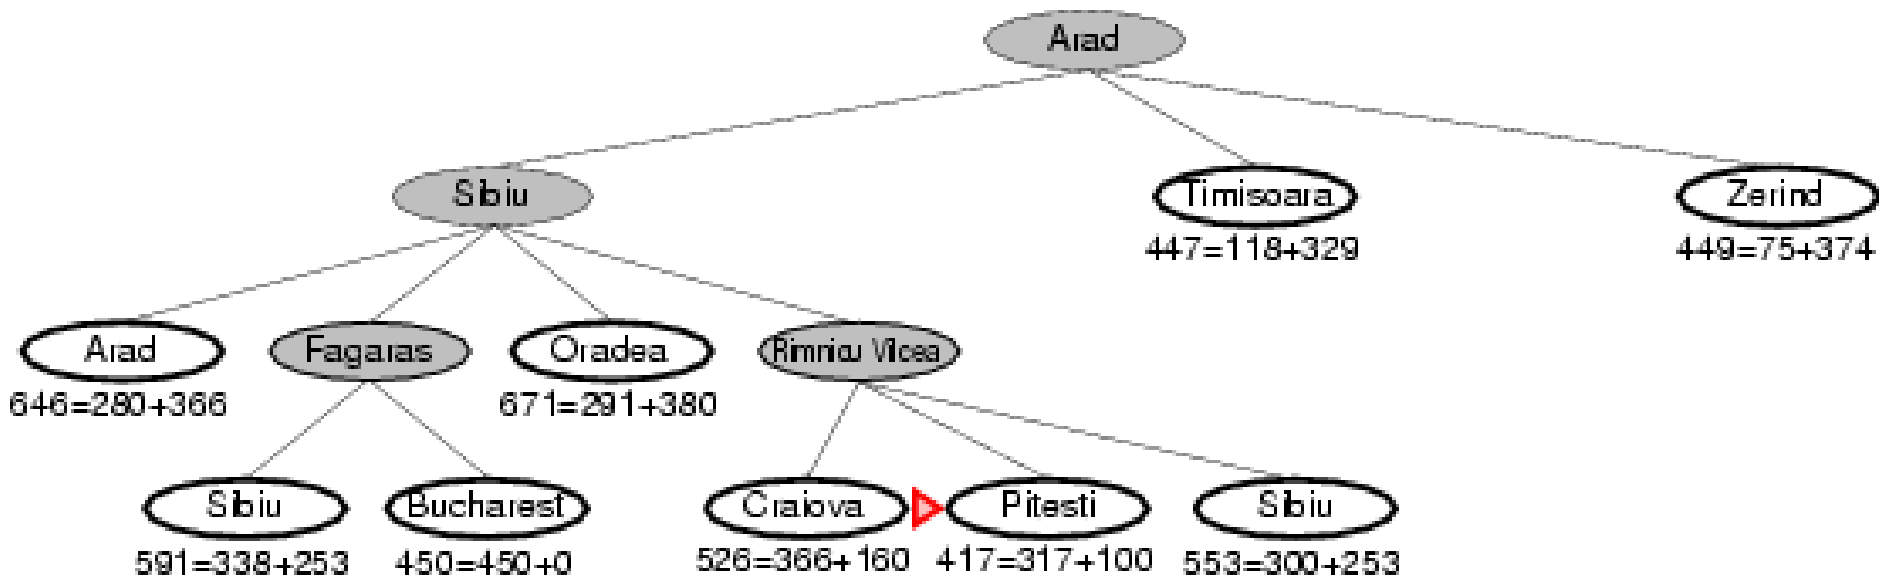
A* Search for Romania Example



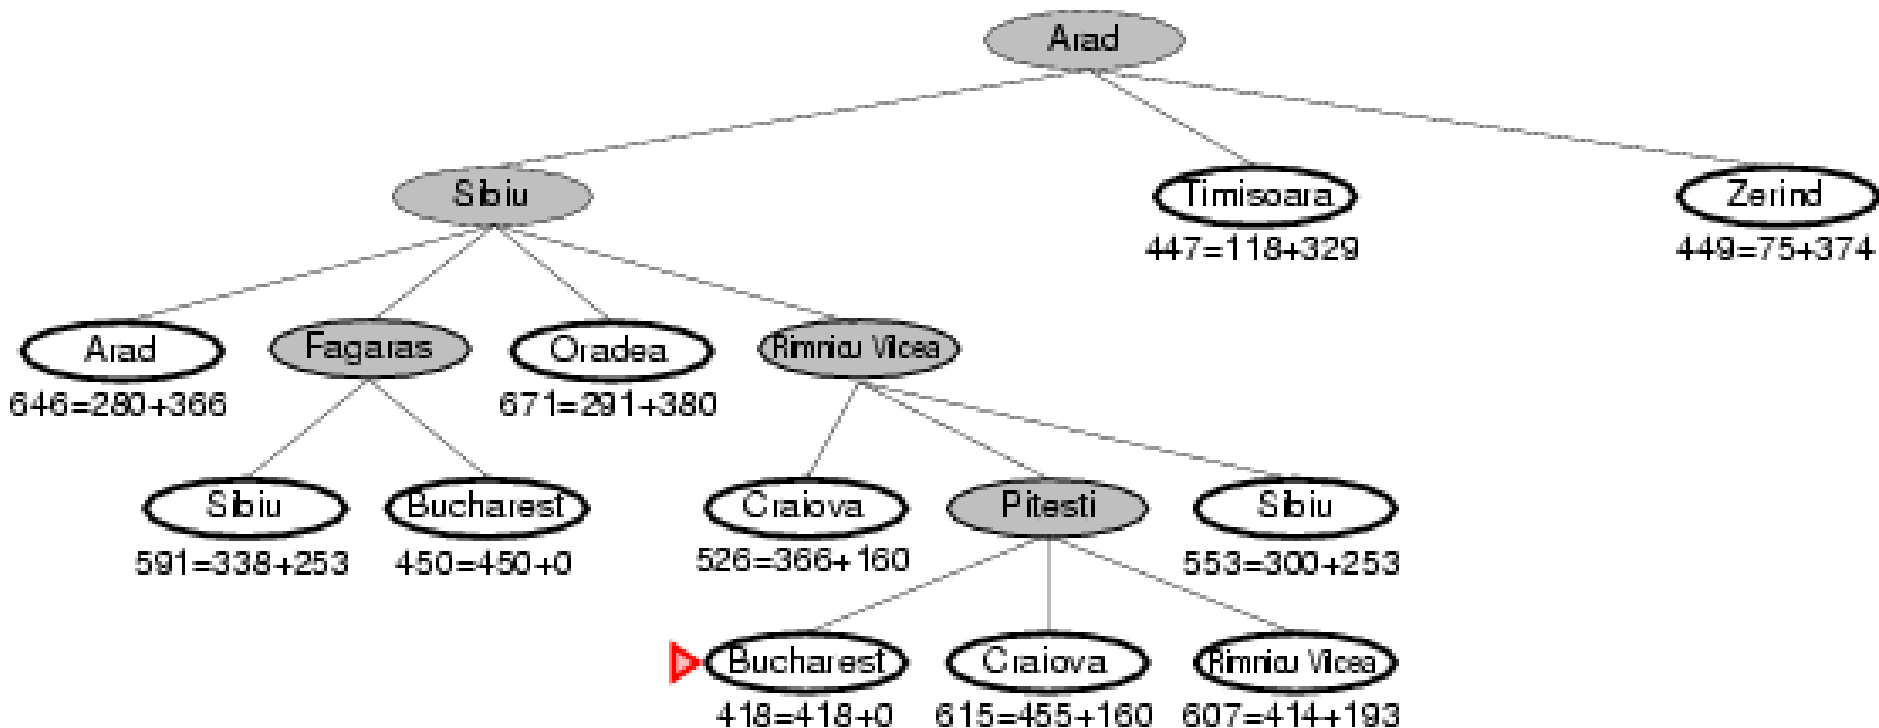
A* Search for Romania Example



A* Search for Romania Example



A* Search for Romania Example




Bucharest is selected for expansion and it is a goal node => stop

Solution path: Arad-Sibiu-Rimnicu Vilcea-Pitesti-Bucharest, cost=418


A* Search – Another Example

- **Given:**
 - **Goal nodes:** C, I, E and K
 - **Step path cost:** along the links
 - ***h* value of each node:** in brackets ()
 - **Same priority nodes -> expand the last added first**
- **Run A***
 - **list of expanded nodes =?**
 - **solution path =?**
 - **cost of the solution:=?**




Solution

- **Fringe:** S
- **Expanded:** nil



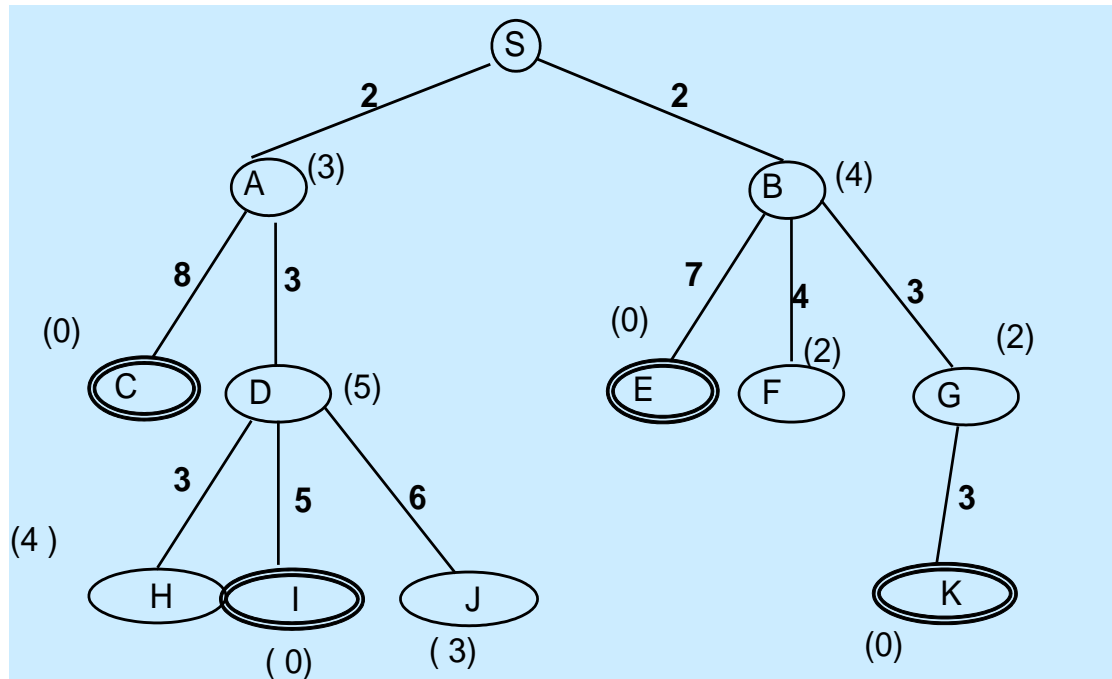
Solution

- **Fringe:** (A, 5), (B, 6) //keep the fringe in sorted order
- **Expanded:** S



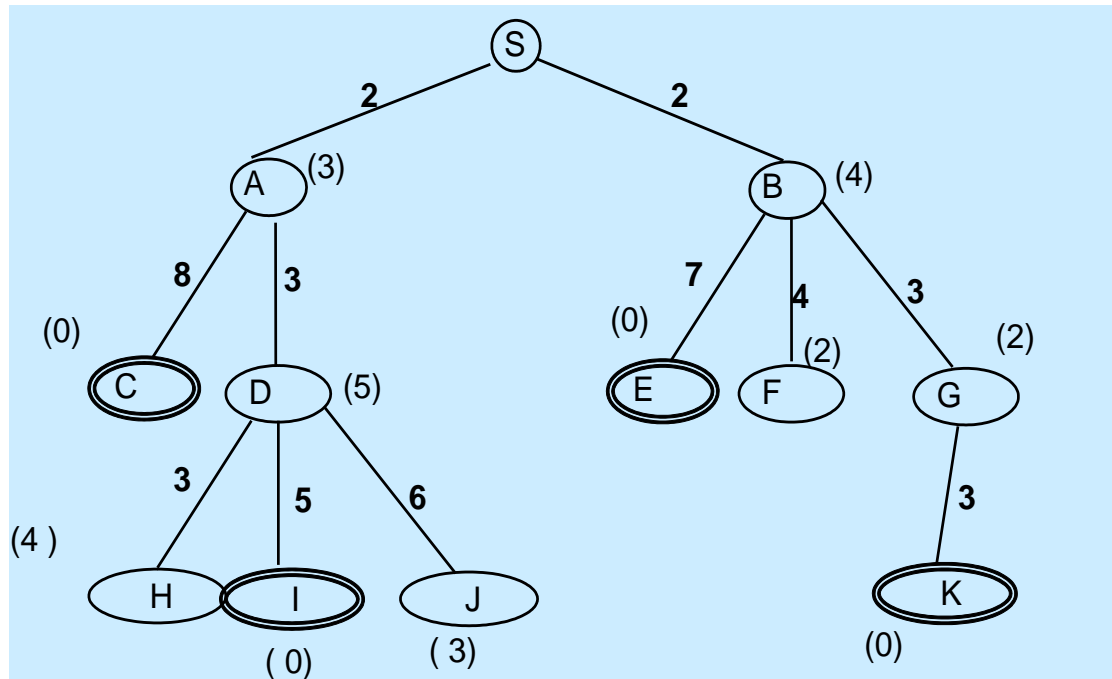
Solution

- **Fringe:** (B, 6), (C, 10), (D, 10) //the added children are in blue
- **Expanded:** S, (A, 5)




Solution

- **Fringe:** (G, 7), (F, 8), (E, 9), (C, 10), (D, 10)
- **Expanded:** S, (A, 5), (B, 6)




Solution

- **Fringe:** (K, 8), (F, 8), (E, 9), (C, 10), (D, 10)
- **Expanded:** S, (A, 5), (B, 6), (G, 7)




Solution

- **K is selected; Goal node? Yes => stop**
- **Expanded: S, A, B, G, K**
- **Solution path: SBGK, cost=8**
- **Is this the optimal solution=?**



A* and UCS

- UCS is a special case of A* when $h(n) = ?$
- In other words, when will A* behave as UCS?




Hint:

- UCS uses which cost?
- A* uses which cost?
- Relation between the 2 costs =?

A* and UCS (Answer)


- UCS is a special case of A* when $h(n) = ?$
- In other words, when will A* behave as UCS?



- UCS uses which cost?
- A* uses which cost?
- UCS: $g(n)$
- A*: $f(n) = g(n) + h(n)$
- if $h(n) = 0 \Rightarrow f(n) = g(n)$,
i.e. A* becomes UCS

A* and BFS


- **BFS is a special case of A* when $f(n) = ?$**
- **When will A* behave as BFS?**



A* and BFS (Answer)

- **BFS is a special case of A* when $f(n) = ?$**



when $f(n) = \text{depth}(n)$



- **And also when this assumption for resolving ties is true: among nodes with the same priority, the left most is expanded first**

BFS, UCS and A*

- BFS is a special case of A* when $f(n)=depth(n)$
- BFS is also a special case of UCS when $g(n)=depth(n)$
- UCS is a special case of A* when $h(n)=0$




Admissible Heuristic

- A heuristic $h(n)$ is admissible if for every node n :
 - $h(n) \leq h^*(n)$ where $h^*(n)$ is the true cost to reach a goal from n
 - i.e. the estimate to reach a goal is smaller than (or equal to) the true cost to reach a goal
- Admissible heuristics are *optimistic* – they think that the cost of solving the problem is less than it actually is!
 - e.g. the straight line distance heuristic $h_{SLD}(n)$ never overestimates the actual road distance (cost from n to goal) => it is admissible
- Theorem: If h is an *admissible heuristic*, then A* is complete and optimal



Is h Admissible for Our Example?

- No need to check goal nodes ($h=0$ for them) and nodes that are not on a goal path
- $h(S)=5 \leq 8$ (shortest path from S to a goal, i.e. to goal K)
- $h(B)=4 \leq 6$
- $h(G)=2 \leq 3$
- $h(A)=3 \leq 8$
- $h(D)=5 \leq 5$
- $\Rightarrow h$ is admissible



Optimality of A* - Proof

Optimal solution = the shortest (lowest cost) path to a goal node


Idea: Suppose that some sub-optimal goal G_2 has been generated and it is in the fringe. We will show that G_2 can not be selected from the fringe.

Given:

G - the optimal goal

G_2 – a sub-optimal goal

h is admissible



To prove: G_2 can not be selected from the fringe for expansion


Proof:

Let n be an unexpanded node in the fringe such that n is on the optimal (shortest) path to G (there must be such a node). We will show that $f(n) < f(G_2)$, i.e. n will be expanded, not G_2

Optimality of A* - Proof (2)

Compare $f(G_2)$ and $f(G)$

- 1) $f(G_2) = g(G_2) + h(G_2)$ (by definition) = $g(G_2)$ as $h(G_2) = 0$, G_2 is a goal
- 2) $f(G) = g(G) + h(G)$ (by definition) = $g(G)$ as $h(G) = 0$, G is a goal
- 3) $g(G_2) > g(G)$ as G_2 is suboptimal
- 4) $\Rightarrow f(G_2) > f(G)$ by substituting 1) and 2) into 3)



Optimality of A* - Proof (3)

Compare $f(n)$ and $f(G)$

5) $f(n) = g(n) + h(n)$ (by definition)

6) $h(n) \leq h^*(n)$ where $h^*(n)$ is the true cost from n to G (as h is admissible)

7) $\Rightarrow f(n) \leq g(n) + h^*(n)$ (5 & 6)

8) $= g(G)$ path cost from S to G via n

9) $g(G) = f(G)$ as $f(G) = g(G) + h(G) = g(G) + 0$ as $h(G) = 0$, G is a goal


10) $\Rightarrow f(n) \leq f(G)$ (7,8,9)

Thus $f(G) < f(G_2)$ (4)

$f(n) \leq f(G)$ (10)

11) $f(n) \leq f(G) < f(G_2)$ (10, 4)

12) $f(n) < f(G_2) \Rightarrow n$ will be expanded not G_2 ; A* will not select G_2 for expansion



Admissible Heuristics for 8-puzzle – h_1

- $h_1(n)$ = number of misplaced tiles
- $h_1(\text{Start}) = ?$
 - 7 (7 of 8 tiles are out of position)
- Why is h_1 admissible?
 - recall: admissible heuristics are optimistic – they never overestimate the number of steps to the goal
 - h_1 : any tile that is out of place must be moved once
 - true cost: higher; any tile that is out of place must be moved at least once

5	4	
6	1	8
7	3	2

Start State

1	2	3
8		4
7	6	5

Goal State

Admissible Heuristics for 8-puzzle – h_2

- $h_2(n)$ = the sum of the distances of the tiles from their goal positions (Manhattan distance)
 - note: tiles can move only horizontally and vertically
- $h_2(Start) = ?$
 - 18 (2+3+3+2+4+2+0+2)
- Why is h_2 admissible?
 - h_2 : at each step move a tile to an adjacent position so that it is 1 step closer to its goal position and you will reach the solution in h_2 steps, e.g. move tile 1 up, then left
 - True cost: higher as moving a tile to an adjacent position is not always possible; depends on the position of the blank tile

5	4	
6	1	8
7	3	2

Start State

1	2	3
8		4
7	6	5

Goal State

Dominance

- Definition of a *dominant heuristic*:
 - Given 2 admissible heuristics h_1 and h_2 ,
 - h_2 *dominates* h_1 if for all nodes n $h_2(n) \geq h_1(n)$
- Theorem: A* using h_2 will expand fewer nodes than A* using h_1 (i.e. h_2 is better for search)
 - $\forall n$ with $f(n) < f^*$ will be expanded (f^* =cost of optimal solution path)
 - $\Rightarrow \forall n$ with $h(n) < f^* - g(n)$ will be expanded
 - but $h_2(n) \geq h_1(n)$
 - $\Rightarrow \forall n$ expanded by A* using h_2 will also be expanded by h_1 and h_1 may also expand other nodes
- Typical search costs for 8-puzzle with $d=14$:
IDS = 3 473 941 nodes, A*(h_1) = 539 nodes, A*(h_2) = 113 nodes
- Dominant heuristics give a better estimate of the true cost to a goal G

Question

- Suppose that $h1$ and $h2$ are two admissible heuristics for a given problem. We define two other heuristics:
 - $h3 = \min(h1, h2)$
 - $h4 = \max(h1, h2)$
- Q1. Is $h3$ admissible?
- Q2. Is $h4$ admissible?
- Q3. Which one is a better heuristic - $h3$ or $h4$?

Answer

- Suppose that $h1$ and $h2$ are two admissible heuristics for a given problem. We define two other heuristics:
 - $h3 = \min(h1, h2)$
 - $h4 = \max(h1, h2)$
- Q1. Is $h3$ admissible?
- Q2. Is $h4$ admissible?
- Q2. Which one is a better heuristic - $h3$ or $h4$?

Answer:

- Q1 and Q2: Both $h3$ and $h4$ are admissible as their values are never greater than an admissible value $h1$ or $h2$
- Q3: $h4$ is a better heuristic since it is closer to the real cost, i.e. $h4$ is a dominant heuristic since $h4(n) \geq h3(n)$

How to Invent Admissible Heuristics?

- By formulating a *relaxed* version of the problem and finding the *exact* solution. This solution is an admissible heuristic.
- Relaxed problem – a problem with fewer restrictions on the actions
- 8-puzzle relaxed formulation 1:
 - a tile can move *anywhere*
 - How many steps do we need to reach the goal state from the initial state? (=solution)
 - solution = the number of misplaced tiles = $h_1(n)$
- 8-puzzle relaxed formulation 2:
 - a tile can move to *any adjacent square*
 - solution = Manhattan distance = $h_2(n)$

Admissible Heuristics from Relaxed Problems

- **Theorem:** The optimal solution to a relaxed problem is an admissible heuristic for the original problem
- Intuitively, this is true because:
The optimal solution to the original problem is also a solution to the relaxed version (by definition) \Rightarrow it must be at least as expensive as the optimal solution to the relaxed version \Rightarrow the solution to the relaxed version is less or equally expensive than the solution to the original problem \Rightarrow it is an admissible heuristic for the original problem

Constructing Relaxed Problems Automatically

- Relaxed problems can be constructed automatically if the problem definition is written in a formal language
 - Problem:
A tile can move from square A to square B if A is adjacent to B and B is blank
 - 3 relaxed problems generated by removing 1 or both conditions:
 - 1) *A tile can move from square A to square B if A is adjacent to B*
 - 2) *A tile can move from square A to square B if B is blank*
 - 3) *A tile can move from square A to square B* (always, no conditions)
- ABSOLVER (1993) is a program that can generate heuristics automatically using the “relaxed problem” method and other methods
 - Generated a new heuristic for the 8-puzzle that was better than any existing heuristic
 - Found the first useful heuristic for the Rubik’s cube puzzle



No Single Clearly Best Heuristic?

- Often we can't find a single heuristic that is clearly the best (i.e. dominant)
- We have a set of heuristics h_1, h_2, \dots, h_m but none of them dominates any of the others
- Which should we choose?
- Solution: define a composite heuristic:
$$h(n) = \max\{h_1(n), h_2(n), \dots, h_m(n)\}$$

At a given node, it uses whichever heuristic is most accurate (dominant)
- Is $h(n)$ admissible?
Yes, because the individual heuristics are admissible

Learning Heuristics from Experience

- **Example: 8-puzzle**
- **Experience = many 8-puzzle solutions (paths from A to B)**
- **Each previous solution provides a set of examples to learn h**
- **Each example is a pair (state, associated h)**
 - h is known for each state, i.e. we have a *labelled* dataset
- **The state is suitably represented as a set of useful features, e.g.**
 - $f1$ = number of misplaced tiles
 - $f2$ = number of adjacent tiles that should not be adjacent
 - h is a function of the features but we don't know how exactly it depends on them, we will learn this relationship from the data
- **We can generate e.g. 100 random 8-puzzle configurations and record the values of $f1$, $f2$ and h to form a *training set* of examples. Using this training set, we build a classifier.**
- **We use this classifier on new data, i.e. given $f1$ and $f2$, to predict h which is unknown. No guarantee that the learned heuristic is admissible or consistent.**


training data

Ex.#	f1	f2	h
Ex1	7	8	14
...			
Ex100	5	2	5

Back to A* and another property of the heuristics...

Consistent (Monotonic) Heuristic




- **Consider a pair of nodes ni and nj , where ni is the parent of nj**



- h is a *consistent (monotonic) heuristic*, if for all such pairs in the search graph the following triangle inequality is satisfied:

$$h(ni) \leq cost(ni,nj) + h(nj) \text{ for all } n$$

parent
child







Another Interpretation of the Triangle Inequality

$$h(n_i) \leq \text{cost}(n_i, n_j) + h(n_j) \text{ for all } n$$

parent

child

- $\Rightarrow h(n_j) \geq h(n_i) - \text{cost}(n_i, n_j)$, i.e. along any path our estimate of the remaining cost to the goal cannot decrease by more than the arc cost



Consistency Theorems

- Theorem 1:** If $h(n)$ is consistent, then $f(n_j) \geq f(n_i)$, i.e. f is non-decreasing along any path

child
parent

Given: $h(n_i) \leq c(n_i, n_j) + h(n_j)$

To prove: $f(n_j) \geq f(n_i)$


Proof: $f(n_j) = g(n_j) + h(n_j) =$

$$= g(n_i) + c(n_i, n_j) + h(n_j) =$$

$$\geq g(n_i) + h(n_i) =$$

$$= f(n_i)$$

$$\Rightarrow f(n_j) \geq f(n_i)$$



deff. $h(n)$ consistent

- Theorem 2:** If $f(n_j) \geq f(n_i)$, i.e. f is non-decreasing along any path, then $h(n)$ is consistent


Admissibility and Consistency

- **Consistency is the stronger condition**
- **Theorems:**
 - **If a heuristic is consistent, it is also admissible**
consistent \Rightarrow admissible
 - **If a heuristic is admissible, there is no guarantee that it is consistent**
admissible \nRightarrow consistent

Completeness of A* with Consistent Heuristic – Intuitive Idea


- A* uses the f-cost to select nodes for expansion
- If h is consistent, the f-costs are non-decreasing \Rightarrow we can draw f-contours in the state space
- A* expands nodes in order of increasing f-values, i.e.
 - It gradually adds f-contours of nodes
 - Nodes inside a contour have f-cost less than or equal to the contour value

- Completeness – as we add bands of increasing f, we must eventually reach a band where $f=h(G)+g(G)=h(G)$



Optimality of A* with Consistent Heuristic – Intuitive Idea

- A* finds the optimal solution, i.e. the one with smallest path cost $g(n)$ among all solutions
- The first solution must be the optimal one, as subsequent contours will have higher f-cost, and thus higher g-cost ($h(n)=0$ for goal nodes):
 - Bands $f_1 < f_2 < f_3 \dots$
 - Compare 2 solutions at band 2 and 3: G2 and G3 (G2 will be found first)
 - $f(G_2) = g(G_2) + h(G_2)$, $f(G_3) = g(G_3) + h(G_3)$
 - But $f(G_2) < f(G_3)$ and $h(G_2) = h(G_3) = 0 \Rightarrow g(G_2) < g(G_3)$, i.e. the first solution found is the optimal



A* with Consistent Heuristic is Optimally Efficient

- **Theorem:** If h is a consistent heuristic, then A* is *optimally efficient* among all optimal search algorithms using h
 - no other optimal algorithm using h is guaranteed to expand fewer nodes than A*
- Which are the optimal algorithms we have studied so far?
- Which are the optimal heuristic algorithms we have studied so far?

Properties of A*

- **Complete?** Yes, unless there are infinitely many nodes with $f \leq f(G)$, G – optimal goal state
- **Optimal?** Yes, with admissible heuristic
- **Time?** Exponential $O(b^d)$
- **Space?** Exponential, keeps all nodes in memory

- For most problems, the number of nodes which have to be expanded is exponential
- Both time and space are problems for A* but space is the bigger problem - A* runs out of space long before it runs out of time; solution: Iterative Deepening A* (IDA*) or Simplified Memory-Bounded A* (SMA*)

Summary of A*

- An admissible heuristic never overestimates the true distance to a goal
- A consistent (monotonic) heuristic satisfies the triangle equation
- $h(n)$ satisfies the triangle equation $\Leftrightarrow f(n)$ does not decrease along any path
- Admissible \nRightarrow consistent
- Consistent \Rightarrow admissible
- Dominant heuristic
 - given 2 admissible heuristics h_1 and h_2 , h_2 is dominant if it gives a better estimate of the true cost to a goal node
 - A* with a dominant heuristic will expand fewer nodes

Summary of A* (2)

- If $h(n)$ is admissible, A* is optimal
- If $h(n)$ is consistent, A* is optimally efficient - A* will expand less or equal number of nodes than any other optimal algorithm using $h(n)$
- However, theoretical completeness and optimality do not mean practical completeness and optimality if it takes too long to get the solution (time and space are exponential)
- \Rightarrow If we can't design an accurate admissible or consistent heuristic, it may be better to settle for a non-admissible heuristic that works well in practice or for a local search algorithm (next lecture) even though completeness and optimality are no longer guaranteed.
- \Rightarrow Also, although dominant (i.e. good) heuristics are better, they may need a lot of time to compute; it may be better to use a simpler heuristic - more nodes will be expanded but overall the search may be faster.