

COMP3308/COMP3608, Lecture 3b

ARTIFICIAL INTELLIGENCE

Local Search Algorithms

Reference: Russell and Norvig, ch. 4

Outline


- **Optimisation problems**
- **Local search algorithms**
 - **Hill-climbing**
 - **Beam search**
 - **Simulated annealing**
 - **Genetic algorithms**

Optimisation Problems

- **Problem setting so far: path finding**
 - **Goal:** find a path from S to G
 - **Solution:** the path
 - **Optimal solution:** least cost path
 - **Search algorithms:**
 - **Uninformed:** BFS, UCS, DFS, IDS
 - **Informed:** greedy, A*
- **Now a new problem setting: optimisation problem**
 - **Each state has a value v**
 - **Goal:** find the **optimal state**
 - = the state with the highest or lowest v score (depending on what is desirable, maximum or minimum)
 - **Solution:** the state; the path is not important
- **A large number of states => can't be enumerated**
 - => **We can't apply the previous algorithms – too expensive**


Optimisation Problems - Example

- **n-queens problem**
 - The solution is the goal configuration, not the path to it
- **Non-incremental formulation**
 - *Initial state*: n-queens on the board (given or randomly chosen)
 - *States*: any configuration with n-queens on the board
 - *Goal*: no queen is attacking each other
 - *Operators*: “move a queen” or “move a queen to reduce the number of attacks”



V-value Landscape

- Each state has a value v that we can compute
- This value is defined by a heuristic *evaluation function* (also called *objective function*)
- Goal - 2 variations depending on the task:
 - find the state with the highest value (global maximum) or
 - find the state with the lowest value (global minimum)
- **Complete** local search – finds a goal state if one exists
- **Optimal** local search – finds the best goal state – the state associated with the global maximum/minimum



Hill-Climbing Algorithm - Idea



- Also called *iterative improvement* algorithm
- Idea: Keep only a single state in memory, try to improve it
- Two variations:
 - Steepest *ascent* – the goal is the *maximum* value
 - Steepest *descent* – the goal is the *minimum* value

Hill-climbing Search

- Assume that we are looking for a maximum value (i.e. hill-climbing ascend)
- Idea: move around trying to find the highest peak
 - Store only the current state
 - Do not look ahead beyond the immediate neighbors of the current state
 - If a neighboring state is better, move to it and continue, otherwise stop
 - “Like climbing Everest in thick fog with amnesia”

we can't see the whole landscape,
only the neighboring states

keeps only 1 state in memory, no
backtracking to previous states



Hill-climbing Algorithm

Hill-climbing descent

1) Set current node n to the initial state s

(The initial state can be given or can be randomly selected)

2) Generate the successors of n . Select the best successor n_{best} ; *it is the successor with the best v score, $v(best)$ (i.e. the lowest score for descent)*


**3) If $v(best) > v(n)$, return n //compare the child with the parent; if child not
//better – stop; local or global minimum found**

**Else set n to n_{best} . Go to step 2 //if better - accept the child and keep
// searching**

- **Summary: Always expands the best successor, no backtracking**

Hill-climbing – Example 1


- v - value is in brackets; the lower, the better (i.e. descent)
- Expanded nodes: SAF



Hill-climbing – Example 2



- Given: 3x3 grid, each cell is colored either red (R) or blue (B)
- Aim: find the coloring with minimum number of pairs of adjacent cells with the same color
- Ascending or descending?

v – # pairs of adjacent cells with the same color



Picture from N. Nielsen, AI, 1998

Hill-climbing Search



Weaknesses:

- Not a very clever algorithm – can easily get stuck in a local optimum (maximum/minimum)
- However, not all local maxima/minima are bad – some may be reasonably good even though not optimal

Advantages: good choice for hard, practical problems

- Uses very little memory
- Finds reasonable solutions in large spaces where systematic algorithms are not useful


Not complete, not optimal but keeps just one node in memory!

Hill-climbing – Escaping Bad Local Optima

- Hill climbing finds the closest local optimum (minimum or maximum, depending on the version – descent or ascent)
 - Which may or may not be the global optimum
- The solution that is found depends on the initial state
- When the solution found is not good enough - random restart:
 - run the algorithm several times starting from different points; select the best solution found (i.e. the best local optimum)
 - *If at first you don't succeed, try, try again!*
 - This is applicable for tasks without a fixed initial state

Hill-climbing – Escaping Bad Local Optima (2)


- Plateaus (flat areas): no change or very small change in v
- Our version of hill-climbing does not allow visiting states with the same v as it terminates if the best child's v is the same as the parent's
- But other versions keep searching if the values are the same and this may result in visiting the same state more than once and walking endlessly
- Solution: keep track of the number of times v is the same and do not allow revisiting of nodes with the same v



Hill-climbing – Escaping Bad Local Optima (3)

- **Ridges** – the current local maximum is not good enough; we would like to move up but all moves go down

- **Example:**



- **Dark circles = states**
- **A sequence of local maxima that are not connected with each other**
- **From each of them all available actions point downhill.**


- **Possible solutions: combine 2 or more moves in a macro move that can increase the height or allow a limited number of look-ahead search**

Beam Search

- It keeps track of k states rather than just 1
- **Version 1: Starts with 1 given state**
 - At each level: generate all successors of the given state
 - If any one is a goal state, stop; else select the k best successors from the list and go to the next level
- **Version 2: Starts with k randomly generated states**
 - At each level: generate all successors of all k states
 - If any one is a goal state, stop; else select the k best successors from the list and go to the next level
- **In nutshell: keeps only k best states**

Beam Search - Example

- Consider the version that starts with 1 given state
- Starting from S, run beam search with $k=2$ using the values in brackets as evaluation function (the smaller, the better)
- Expanded nodes = ?



- S
- generate ABC
- select AB (the best 2 children)
- generate DEFH
- select FH (the best 2 children)
- expanded nodes: SABFH

Beam Search and Hill-Climbing Search

- **Compare beam search with 1 initial state and hill climbing with 1 initial state**
 - **Beam – 1 start node, at each step keeps k best nodes**
 - **Hill climbing – 1 start node, at each step keeps 1 best node**
- **Compare beam search with k random initial states and hill-climbing with k random initial states**
 - **Beam – k starting positions, k threads run in parallel, passing of useful information among them as at each step the best k children are selected**
 - **Hill climbing – k starting positions, k threads run individually, no passing of information among them**


Beam Search with A*

- **Recall that memory was a big problem for A***
- **Idea: keep only the best k nodes in the fringe, i.e. use a priority queue of size k**
- **Advantage: memory efficient**
- **Disadvantage: neither complete, nor optimal**

Simulated Annealing

- **What is annealing in metallurgy?**
 - a material's temperature is gradually decreased (very slowly) allowing its crystal structure to reach a minimum energy state
 - **Similar to hill-climbing but selects a random successor instead of the best successor (step 2 below)**
-
- 1) **Set current node n to the initial state s .**
Randomly select m , one of n 's successors
 - 2) **If $v(m)$ is better than $v(n)$, $n=m$ //accept the child m**
Else $n=m$ with a probability p //accept the child m with probability p
 - 3) **Go to step 2 until a predefined number of iterations is reached or the state reached (i.e. the solution found) is good enough**

The Probability p



- Assume that we are looking for a minimum
- There are different ways to define p , e.g.
- Main ideas:
 - 1) p decreases exponentially with the badness of the child (move) and
 - 2) bad children (moves) are more likely to be allowed at the beginning than at the end
- nominator – shows how good the child m is
 - Bad move (the child is worse than the parent): $v(n) < v(m)$, e.g.
 - case1: $v(n)=10, v(m)=20, p1=e^{-10/T}$
 - case2: $v(n)=10, v(m)=11, p2=e^{-1/T}$
 - m (the child) in case1 is worse than in case2
 - $p1 < p2$ as T is positive $\Rightarrow p$ exponentially decreases with the badness of the move

$$p = e^{\frac{v(n)-v(m)}{T}}$$

The Probability p (2)

$$p = e^{\frac{v(n)-v(m)}{T}}$$

- **denominator: parameter T that decreases (anneals) over time based on a *schedule*, e.g. $T=T*0.8$**
 - **high T – bad moves are more likely to be allowed**
 - **low T – more unlikely; becomes more like hill-climbing**
 - **T decreases with time and depends on the number of iterations completed, i.e. until “bored”**
- **Some versions have an additional step (Metropolis criterion for accepting the child):**
 - **p is compared with p' , a randomly generated number $[0,1]$**
 - **If $p > p'$, accept the child, otherwise reject it**
- **In summary, simulated annealing combines a **hill-climbing step** (accepting the best child) with a **random walk step** (accepting bad children with some probability). The random walk step can help escape bad local minima.**


Simulated Annealing - Theorem

- What is the correspondence?
 - v – total energy of the atoms in the material
 - T – temperature
 - *schedule* – the rate at which T is lowered
- **Theorem:** If the schedule lowers T slowly enough, the algorithm will find global optimum
 - i.e. is complete and optimal given a long enough cooling schedule => annealing schedule is very important
- Simulated annealing has been widely used to solve VLSI layout problems, factory scheduling and other large-scale optimizations
- It is easy to implement but a “slow enough” schedule may be difficult to set

Genetic Algorithms

- Inspired by mechanisms used in evolutionary biology, e.g. selection, crossover, mutation
- Similar to beam search, in fact a variant of stochastic beam search
- Each state is called an *individual*. It is coded as a string.
- Each state n has a fitness score $f(n)$ (evaluation function). The higher the value, the better the state.
- Goal: starting with k randomly generated individuals, find the optimal state
- Successors are produced by **selection, crossover and mutation**
- At any time keep a fixed number of states (the population)

Example – 8-queens Problem



One possible encoding is (3 2 7 5 2 4 1 1)

column 1: a queen
at position 3

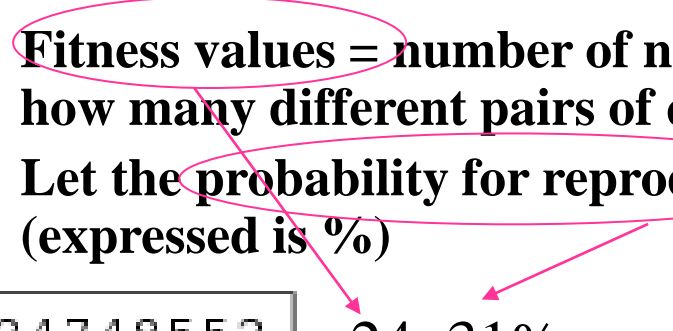
column 2: a queen
at position 2

...

Example – 8-queens Problem (2)

• Suppose that we are given 4 individuals (initial population) with their fitness values

- **Fitness values = number of non-attacking pairs of queens** (given 8 queens, how many different pairs of queens are there? $28 \Rightarrow$ max value is 28)
- **Let the probability for reproduction is proportional to the fitness** (expressed in %)




24748552	24	31%
32752411	23	29%
24415124	20	26%
32543213	11	14%

(a)


Initial Population

Example – 8-queens Problem (3)

- **Select** 4 individuals for reproduction based on the fitness function
 - The higher the fitness function, the higher the probability to be selected
 - Let individuals 2, 1, 2 and 3 be selected, i.e. individual 2 is selected twice while 4 is not selected
- **Crossover** – random selection of crossover point; crossing over the parents strings
- **Mutation** – random change of bits (in this case 1 bit was changed in each individual)



A Closer Look at the Crossover



$$(3 \text{ } 2 \text{ } 7 \text{ } 5 \text{ } 2 \text{ } 4 \text{ } 1 \text{ } 1) + (2 \text{ } 4 \text{ } 7 \text{ } 4 \text{ } 8 \text{ } 5 \text{ } 5 \text{ } 2) = (3 \text{ } 2 \text{ } 7 \text{ } 4 \text{ } 8 \text{ } 5 \text{ } 5 \text{ } 2)$$

- When the 2 states are different, crossover produces a state which is a long way from either parents
- Given that the population is diverse at the beginning of the search, crossover takes **big** steps in the state space early in the process and **smaller** later, when more individuals are similar

Genetic Algorithm – Pseudo Code (1 variant)

from <http://pages.cs.wisc.edu/~jerryzhu/cs540.html>

1. Let s_1, \dots, s_N be the current population
 2. Let $p_i = f(s_i) / \sum_j f(s_j)$ be the reproduction probs
 3. FOR $k = 1; k < N; k += 2$
 - parent1 = randomly pick s with probs p
 - parent2 = randomly pick another s with probs p
 - randomly select a crossover point, swap strings of parents 1, 2 to generate children $t[k], t[k+1]$
 4. FOR $k = 1; k \leq N; k++$
 - Randomly mutate each position in $t[k]$ with a small probability
 5. The new generation replaces the old: $\{s\} \leftarrow \{t\}$.
- Repeat until some individual is fit enough or a predefined maximum number of iterations has been reached**

Genetic Algorithms - Discussion

- **Combine:**
 - **uphill tendency (based on the fitness function)**
 - **random exploration (based on crossover and mutation)**
- **Exchange information among parallel threads - the population consists of several individuals**
- **The main advantage comes from crossover**
- **Success depends on the representation (encoding)**
- **Easy to implement**
- **Not complete, not optimal**

Links

Simulated annealing as a training algorithm for backpropagation neural networks:

- **R.S. Sexton, R.E. Dorsey, J.D. Johnson, *Beyond backpropagation: using simulated annealing for training neural networks*, people.missouristate.edu/randallsexton/sabp.pdf**